


Introduction to MongoDB

Requirements of Modern Application

- Big Data
 - Lots of storage
 - Lots of access
- Ever Changing Features
 - Simpler data models
- Flexible deployments
 - Cloud-enabled

Relational Databases



Persistence

Reporting

Integration


Transactions

SQL

Relational Database

- Impedance Mismatch

The age of Big Data




Not a Relational Database

Google → Bigtable

Amazon → Dynamo

NoSQL



1980

1990

2000

2010

Non-relational

Open-source

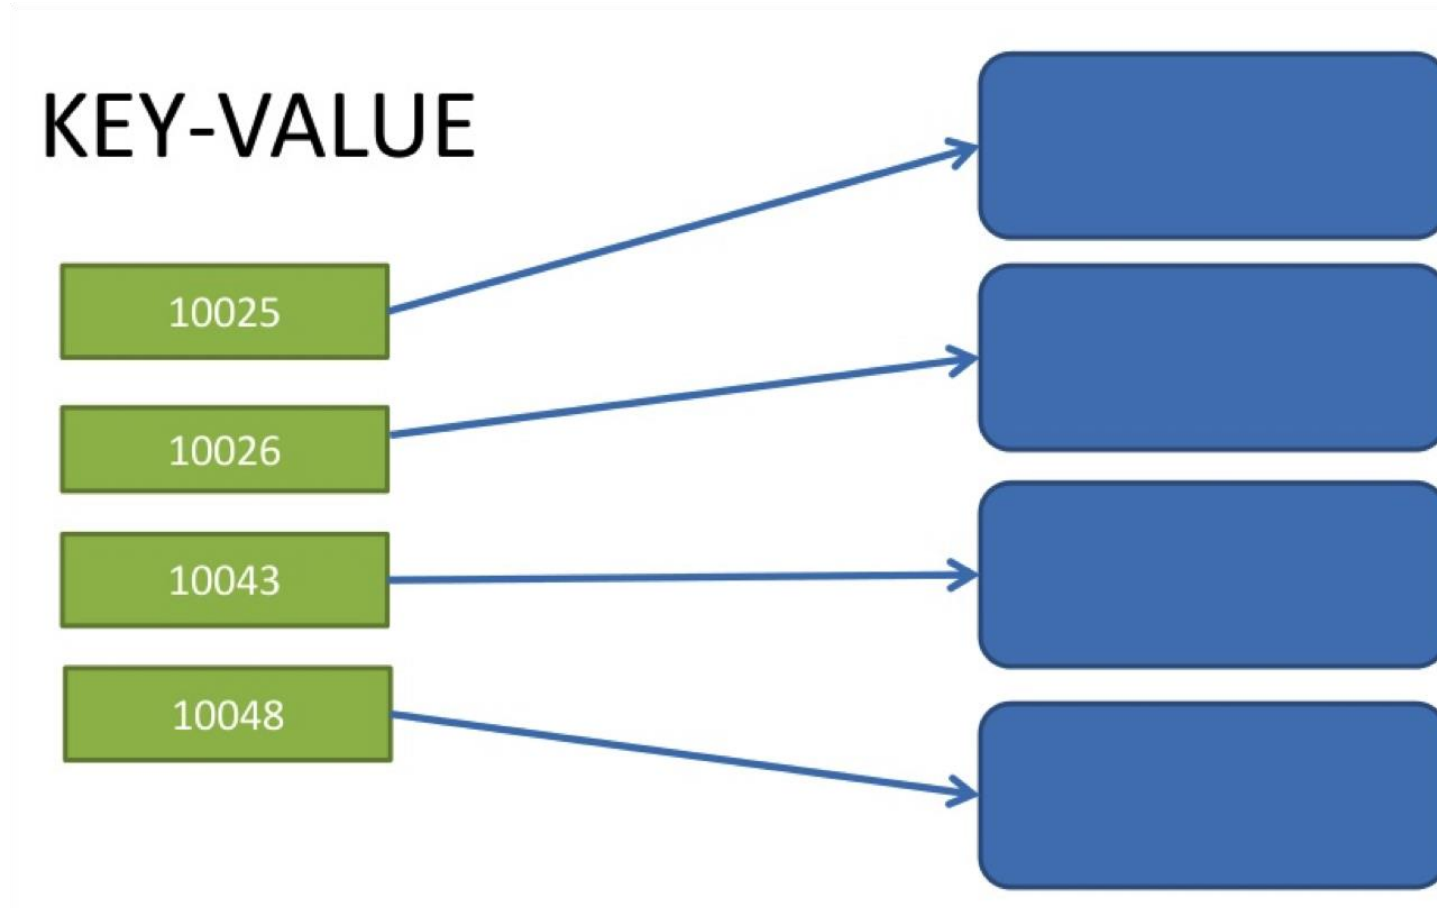
Cluster friendly

Schema-less

Data Model

- Key-value
- Document
- Graph
- Column-family

Key-value




Document

```
{
  "_id": "4c6cd19abac7061798000002",
  "CityId": 42231,
  "CountryID": 1,
  "RegionID": 833,
  "City": "Herat",
  "Latitude": 34.3330001831055,
  "Longitude": 62.2000007629395,
  "TimeZone": "+04:30",
  "DmaId": 0,
  "County": "HERA",
  "Code": ""
}
```

No
Schema

KEY-VALUE



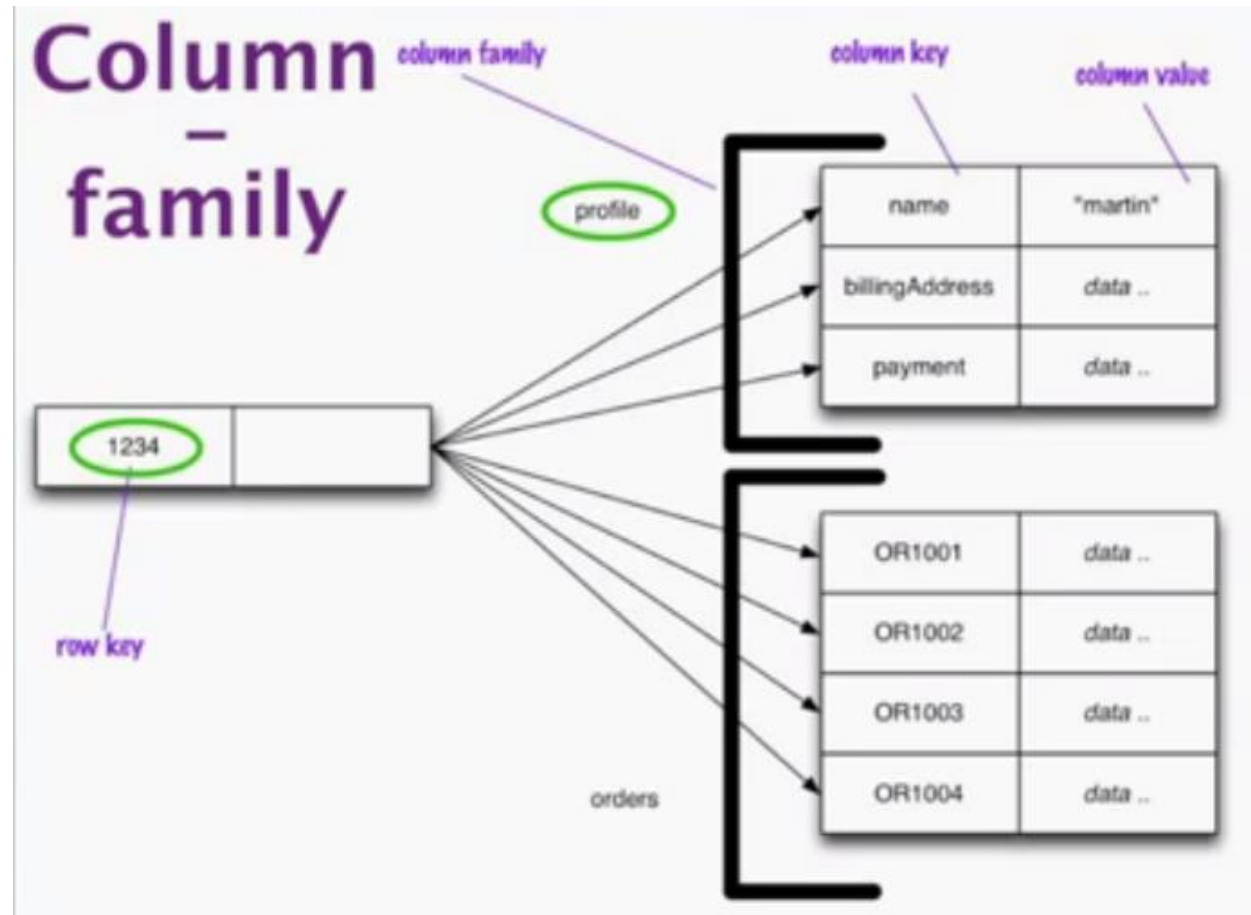
Document



Aggregate

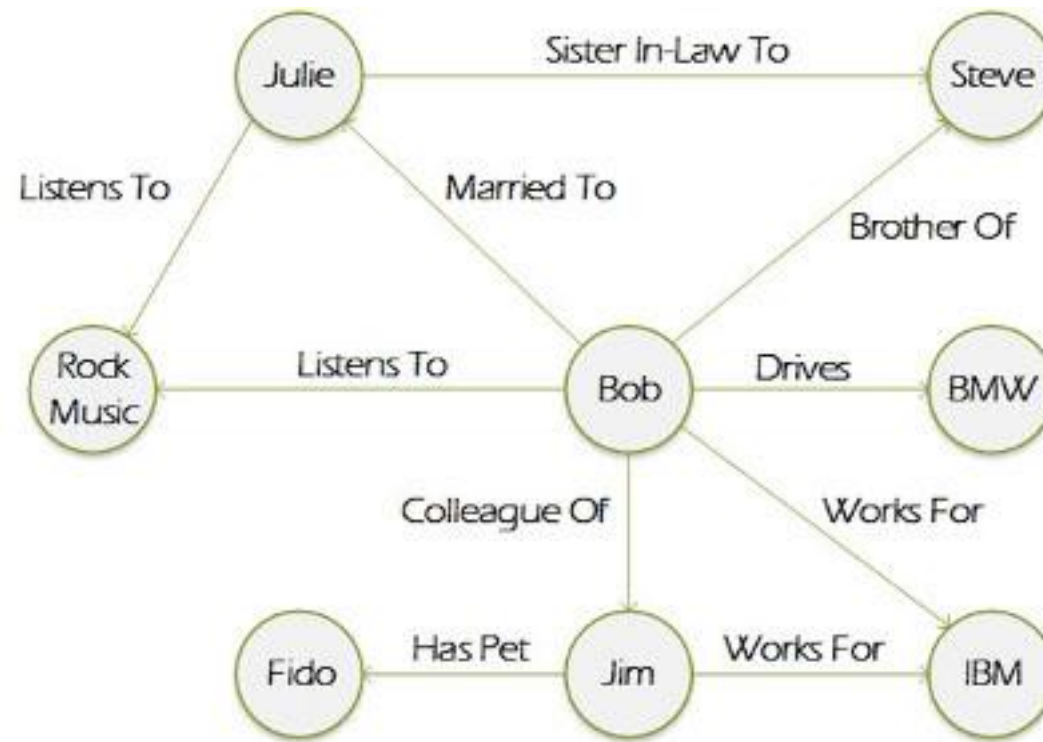
- Aggregate-oriented databse
 - Value
 - Document
- Aggregate is the single unit that we access

Column-family



Aggregate model on Clusters

Graph



NoSQL and Consistency

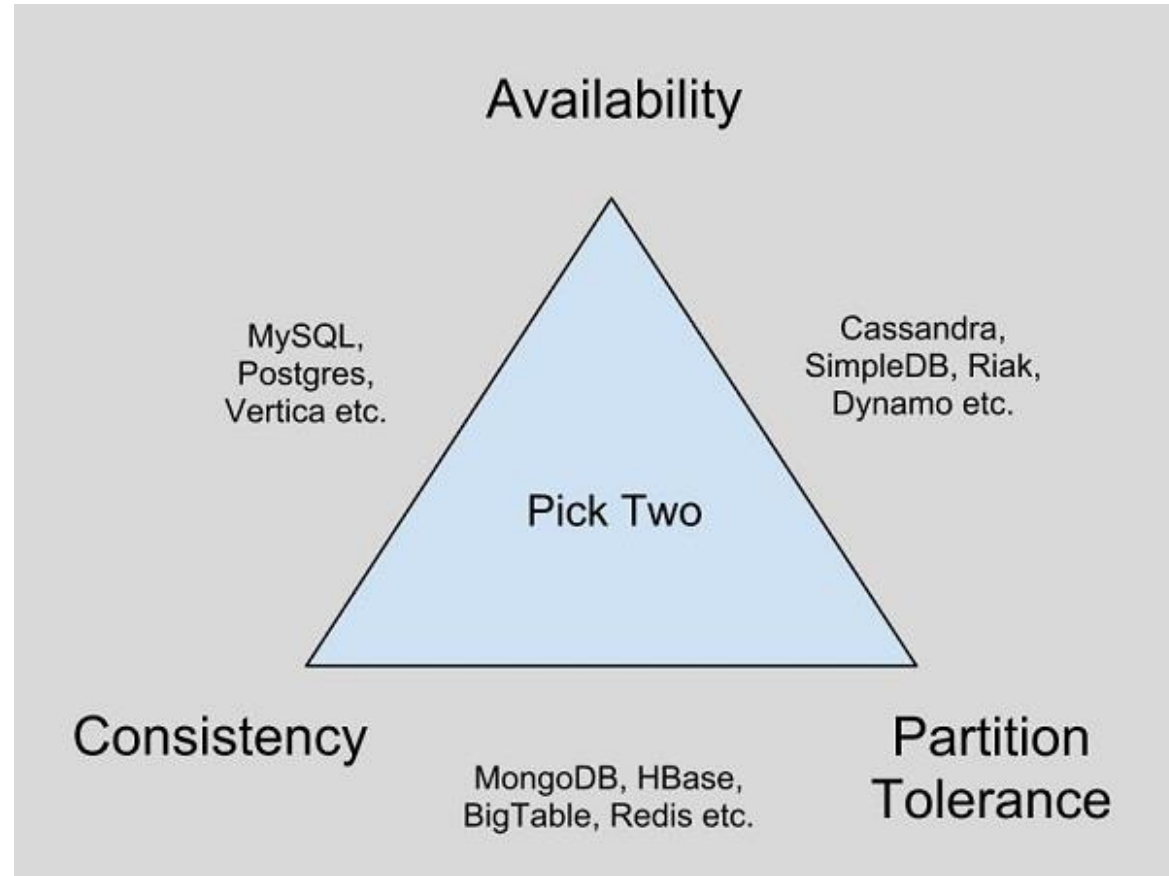
- ACID?

NoSQL and Consistency

- Graph
- Aggregation-oriented
 - Aggregate is the boundary
- Version stamps

NoSQL and Consistency

- Logical
- Replication



Trade-offs

- Consistency vs. Response-time

When and Why to use a NoSQL database

- Large scale data
 - Mobile Cloud Applications
- Easier Development
 - Natural Aggregate
- Web Services/Applications enabling access to data

Introduction

- A NO-SQL database
- Conceptualized to be a distributed database
- Easy to scale out
- Allows schema-free storage of data as collections
- Current version 3.2

Installation

- Windows install
 - Download the Windows installer
 - Create a C:\data\db directory
 - bin\mongod.exe
 - --dbpath PATH to specify any other directory as the base directory
 - bin\mongo.exe
 - Can also be installed as a service

Documents & Collections

- A *document* is the basic unit of data for MongoDB, roughly equivalent to a row in a relational database management system
- A *collection* can be thought of as the schema-free equivalent of a table.

Documents

- Documents: an ordered set of keys with associated values
 - Naturally fits with data structures like map, hash etc.
 - E.g. {"greeting" : "Hello", "foo": 3}
- Key/value pairs in documents are ordered
 - {"greeting" : "Hello", "foo": 3} & {"foo": 3, "greeting" : "Hello"}
 - Values could be several different data types including embedded documents

Keys in Documents

- The keys in a document are strings. Any UTF-8 character is allowed in a key, with a few notable exceptions:
 - Keys must not contain the character `\0` (the null character). This character is used to signify the end of a key.
 - The `.` and `$` characters have some special properties and should be used only in certain circumstances, as described in later chapters. In general, they should be considered reserved.
 - Keys starting with `_` should be considered reserved;
- MongoDB is type-sensitive and case-sensitive. For example, these documents are distinct:
 - `{"foo" : 3}` & `{"foo" : "3"}`
 - `{"foo" : 3}` & `{"Foo" : 3}`
- No duplicate keys
 - `{"greeting" : "Hello, world!"}`, `{"greeting" : "Hello, MongoDB!"}`

Example

```
{
  "_id" : ObjectId("54c955492b7c8eb21818bd09"),
  "address" : {
    "street" : "2 Avenue",
    "zipcode" : "10075",
    "building" : "1480",
    "coord" : [ -73.9557413, 40.7720266 ],
  },
  "borough" : "Manhattan",
  "cuisine" : "Italian",
  "name" : "Vella",
  "restaurant_id" : "41704620"
}
```

source: <https://docs.mongodb.org/getting-started/shell/introduction/>

Collections

- A group of documents like tables are a group of rows
- Collections are schema-free: documents within a single collection can have any number of different “shapes.”
 - {"greeting" : "Hello, world!"}
 - {"foo" : 5}

Collections

- Why do we need separate collections at all?
- Keeping different kinds of documents in the same collection can be a nightmare for developers and admins.
 - Each query should return document of a certain kind
- It is much faster to get a list of collections than to extract a list of the types in a collection.
- Grouping documents of the same kind together in the same collection allows for data locality
- Putting only documents of a single type into the same collection, we can index our collections more efficiently.

Collections - Naming

- The empty string ("") is not a valid collection name.
- Collection names may not contain the character `\0` (the null character.
- You should not create any collections that start with *system.* - a prefix reserved for system collections.
 - For example, the *system.users* collection contains the database's users, and the *system.namespaces* collection contains information about all of the database's collections.
- User-created collections should not contain the reserved character `$` in the name.

Databases

- MongoDB groups collections into databases
- A single instance of MongoDB can host several databases – each of which is completely independent
- Each database is stored in separate files on disk

Databases - Naming

- The empty string ("") is not a valid database name.
- A database name cannot contain any of these characters: ' ' (a single space), ., \$, /, \, or \0 (the null character).
- Database names should be all lowercase.
- Database names are limited to a maximum of 64 bytes.

Reserved Databases

- *admin*
 - This is the “root” database, in terms of authentication. If a user is added to the *admin* database, the user automatically inherits permissions for all databases. There are also certain server-wide commands that can be run only from the *admin* database, such as listing all of the databases or shutting down the server.
- *local*
 - This database will never be replicated and can be used to store any collections that should be local to a single server
- *config*
 - When Mongo is being used in a sharded setup, the *config* database is used internally to store information about the shards.

Mongo Shell

- The shell can connect to daemon running on any other machine
 - bin/mongo [www.server.com:port](#)
 - By default starts with “test” database
 - db variable points to the current database
 - Can use connect to connect with multiple databases
 - bin/mongo localhost:27017/admin – now connects to admin database
 - bin/mongo –nodb - now does not connect with any database

Binary JSON (BSON)

- Binary JSON (BSON): representation of documents that is shared by all drivers, tools, and processes in the MongoDB ecosystem.
- BSON is a lightweight binary format capable of representing any MongoDB document as a string of bytes.
- The database understands BSON, and BSON is the format in which documents are saved to disk.

BSON

- **Lightweight**
 - Overhead is minimum
 - Good to use over the network.
- **Traversable**
 - BSON is designed to be traversed easily.
- **Efficient**
 - Encoding data to BSON and decoding from BSON can be performed very quickly in most languages due to the use of C data types.

Wire Protocol

- Clients communicate with the MongoDB server using a lightweight TCP/IP Wire protocol; default port is 27017
 - A regular TCP/IP Socket
- Two types of messages
 - Client request
 - Database responses

Wire protocol

Standard Message Header

```
struct MsgHeader {  
    int32  messageLength; // total message size, including this  
    int32  requestID;    // identifier for this message  
    int32  responseTo;   // requestID from the original request  
                        // (used in reponses from db)  
    int32  opCode;       // request type - see table below  
}
```

Wire Protocol -opcodes

Opcode Name	opCode value	Comment
OP_REPLY	1	Reply to a client request. responseTo is set
OP_MSG	1000	generic msg command followed by a string
OP_UPDATE	2001	update document
OP_INSERT	2002	insert new document
RESERVED	2003	formerly used for OP_GET_BY_OID
OP_QUERY	2004	query a collection
OP_GET_MORE	2005	Get more data from a query. See Cursors
OP_DELETE	2006	Delete documents
OP_KILL_CURSORS	2007	Tell database client is done with a cursor

References

- https://www.youtube.com/watch?v=ql_g07C_Q5I