

# Overview of MongoDB

# Introduction

- A NO-SQL database
- Conceptualized to be a distributed database
- Easy to scale out
- Allows schema-free storage of data as collections
- Current version 3.2

# Installation

- Windows install
  - Download the Windows installer
  - Create a C:\data\db directory
  - bin\mongod.exe
    - --dbpath PATH to specify any other directory as the base directory
  - bin\mongo.exe
  - Can also be installed as a service

# Wire Protocol

- Clients communicate with the MongoDB server using a lightweight TCP/IP Wire protocol; default port is 27017
  - A regular TCP/IP Socket
- Two types of messages
  - Client request
  - Database responses

# Wire protocol

## Standard Message Header

```
struct MsgHeader {  
    int32  messageLength; // total message size, including this  
    int32  requestID;    // identifier for this message  
    int32  responseTo;   // requestID from the original request  
                        // (used in reponses from db)  
    int32  opCode;       // request type - see table below  
}
```

# Wire Protocol -opcodes

Opcode Name	opCode value	Comment
OP_REPLY	1	Reply to a client request. responseTo is set
OP_MSG	1000	generic msg command followed by a string
OP_UPDATE	2001	update document
OP_INSERT	2002	insert new document
RESERVED	2003	formerly used for OP_GET_BY_OID
OP_QUERY	2004	query a collection
OP_GET_MORE	2005	Get more data from a query. See Cursors
OP_DELETE	2006	Delete documents
OP_KILL_CURSORS	2007	Tell database client is done with a cursor

# Mongo Shell

- The shell can connect to daemon running on any other machine
  - bin/mongo [www.server.com:port](#)
  - By default starts with “test” database
  - db variable points to the current database
    - Can use connect to connect with multiple databases
  - bin/mongo localhost:27017/admin – now connects to admin database
  - bin/mongo –nodb - now does not connect with any database

# Binary JSON (BSON)

- Binary JSON (BSON): representation of documents that is shared by all drivers, tools, and processes in the MongoDB ecosystem.
- BSON is a lightweight binary format capable of representing any MongoDB document as a string of bytes.
- The database understands BSON, and BSON is the format in which documents are saved to disk.



# BSON

- **Lightweight**
  - Overhead is minimum
  - Good to use over the network.
- **Traversable**
  - BSON is designed to be traversed easily.
- **Efficient**
  - Encoding data to BSON and decoding from BSON can be performed very quickly in most languages due to the use of C data types.

# Documents & Collections

- A *document* is the basic unit of data for MongoDB, roughly equivalent to a row in a relational database management system
- A *collection* can be thought of as the schema-free equivalent of a table.

# Documents

- Documents: an ordered set of keys with associated values
  - Naturally fits with data structures like map, hash etc.
  - E.g.     {"greeting" : "Hello", "foo": 3}
- Key/value pairs in documents are ordered
  - {"greeting" : "Hello", "foo": 3}             &     {"foo": 3, "greeting" : "Hello"}
  - Values could be several different data types including embedded documents

# Keys in Documents

- The keys in a document are strings. Any UTF-8 character is allowed in a key, with a few notable exceptions:
  - Keys must not contain the character `\0` (the null character). This character is used to signify the end of a key.
  - The `.` and `$` characters have some special properties and should be used only in certain circumstances, as described in later chapters. In general, they should be considered reserved.
  - Keys starting with `_` should be considered reserved;
- MongoDB is type-sensitive and case-sensitive. For example, these documents are distinct:
  - `{"foo" : 3}` & `{"foo" : "3"}`
  - `{"foo" : 3}` & `{"Foo" : 3}`
- No duplicate keys
  - `{"greeting" : "Hello, world!"}`, `{"greeting" : "Hello, MongoDB!"}`

# Example

```
{
  "_id" : ObjectId("54c955492b7c8eb21818bd09"),
  "address" : {
    "street" : "2 Avenue",
    "zipcode" : "10075",
    "building" : "1480",
    "coord" : [ -73.9557413, 40.7720266 ],
  },
  "area" : "Manhattan",
  "cuisine" : "Italian",
  "name" : "Vella",
  "restaurant_id" : "41704620"
}
```

source: <https://docs.mongodb.org/getting-started/shell/introduction/>

# Collections

- A group of documents like tables are a group of rows
- Collections are schema-free: documents within a single collection can have any number of different “shapes.”
  - {"greeting" : "Hello, world!"}
  - {"foo" : 5}

# Collections

- Why do we need separate collections at all?
- Keeping different kinds of documents in the same collection can be a nightmare for developers and admins.
  - Each query should return document of a certain kind
- It is much faster to get a list of collections than to extract a list of the types in a collection.
- Grouping documents of the same kind together in the same collection allows for data locality
- Putting only documents of a single type into the same collection, we can index our collections more efficiently.

# Collections - Naming

- The empty string ("" ) is not a valid collection name.
- Collection names may not contain the character `\0` (the null character.
- You should not create any collections that start with *system.* - a prefix reserved for system collections.
  - For example, the *system.users* collection contains the database's users, and the *system.namespaces* collection contains information about all of the database's collections.
- User-created collections should not contain the reserved character `$` in the name.



# Databases

- MongoDB groups collections into databases
- A single instance of MongoDB can host several databases – each of which is completely independent
- Each database is stored in separate files on disk

# Databases - Naming

- The empty string ("" ) is not a valid database name.
- A database name cannot contain any of these characters: ' ' (a single space), ., \$, /, \, or \0 (the null character).
- Database names should be all lowercase.
- Database names are limited to a maximum of 64 bytes.

# Reserved Databases

- *admin*
  - This is the “root” database, in terms of authentication. If a user is added to the *admin* database, the user automatically inherits permissions for all databases. There are also certain server-wide commands that can be run only from the *admin* database, such as listing all of the databases or shutting down the server.
- *local*
  - This database will never be replicated and can be used to store any collections that should be local to a single server
- *config*
  - When Mongo is being used in a sharded setup, the *config* database is used internally to store information about the shards.

# Basic Data Types

- Documents are “JSON-like”
  - JSON has only six data types – null, boolean, numeric, string, array, object
- MongoDB has several additional types
  - null, Boolean, string, regular expressions
  - number: default is 64 bit floating, use `NumberInt` or `NumberLong` for 4 byte or 8 byte Integer
    - `{ "x" : NumberInt("3") }`
    - `{ "x" : NumberLong("3") }`
  - date: `{ "x" : new Date() }`
  - embedded documents: `{ "x" : { "foo" : "bar" } }`
  - object id: a 12 byte ID for documents: `{ "x" : ObjectId() }`
  - binary data and code `{ "x" : function() { /* ... */ } }`

# Basic Data Types

- Arrays
  - {"things" : ["pie", 3.14]}
  - arrays can contain different data types as values
  - MongoDB “understands” their structure and knows how to reach inside of arrays to perform operations on their contents.
    - allows us to query on arrays and build indexes using their contents
    - e.g. MongoDB can query for all documents where 3.14 is an element of the "things" array
- Embedded Documents
  - Documents can be used as the *value* for a key.
  - ```
{  
  "name" : "John Doe",  
  "address" : {  
    "street" : "123 Park Street",  
    "city" : "Anytown",  
    "state" : "NY"  
  }  
}
```
  - Same functionality as Arrays

# \_id

- Every document stored in MongoDB must have an "\_id" key.
- The "\_id" key's value can be any type, but it defaults to an ObjectId.
- In a single collection, every document must have a unique value for "\_id", which ensures that every document in a collection can be uniquely identified.

# ObjectIds

- ObjectId is the default type for "\_id".
- The ObjectId class is designed to be lightweight, while still being easy to generate in a globally unique way across different machines.
  - supports distributed nature of MongoDB
- Uses 12 bytes of storage – displayed as 24 digits in hexadecimal

|           |         |     |           |
|-----------|---------|-----|-----------|
| 0 1 2 3   | 4 5 6   | 7 8 | 9 10 11   |
| Timestamp | Machine | PID | Increment |

- if not supplied the \_id is generated at the client end

# Shell

- Shell can run Javascripts
  - can be used to set helper functions, variables etc.
  - mongorc.js is run by the shell at the startup
    - could be used to greet, set global variables, customize prompt etc.



# Creating Documents

- Inserting and Saving Documents
  - `>db.foo.insert({"bar" : "baz"})`
  - return `nInserted` specifying the number of documents inserted.
  - document should be less than 16 MB
    - `>for(var i=0; i<100; i++){db.foo.insert({"foo":"bar", "b":i})}`

# Creating Documents

- Bulk inserts

```
var bulk = db.foo.initializeUnorderedBulkOp();  
bulk.insert( { item: "abc123", defaultQty: 100, status: "A", points: 100 } );  
bulk.insert( { item: "ijk123", defaultQty: 200, status: "A", points: 200 } );  
bulk.insert( { item: "mop123", defaultQty: 0, status: "P", points: 0 } );  
bulk.execute();
```

- allows multiple documents to be inserted at ones
- only useful if you are inserting multiple documents into a single collection
  - you cannot use batch inserts to insert into multiple collections with a single request
- Message size should be less than 48 MB

# Bulk Insert

- Bulk operations builder used to construct a list of write operations to perform in bulk for a single collection
- To initiate the builder  
`db.collection.initializeOrderedBulkOp()` `db.collection.initializeUnorderedBulkOp()` method.
- Ordered Operations
  - MongoDB executes the write operations in the list serially.
    - If an error occurs during the processing of one of the write operations, MongoDB will return without processing any remaining write operations in the list.
- Unordered Operations
  - MongoDB can execute in parallel, as well as in a nondeterministic order.
    - If an error occurs during the processing of one of the write operations, MongoDB will continue to process remaining write operations in the list.

# Removing Documents

- `>db.foo.remove({})`
  - removes all the documents in the foo collection
  - Collection is not removed and meta information also exists
- `>db.foo.remove({"a":"b"})`
  - takes a query document as a parameter
  - only matching documents are removed
- `drop()`: to drop a collection or database
  - once dropped, all is deleted
  - very fast

# Updating Documents

- `update()` is used to change the document
- update parameters
  - an update conditions document to match the documents to update,
  - an update operations document to specify the modification to perform, and
  - an options document
- To change a field value, MongoDB provides update operators, such as `$set` to modify values.
  - `>db.foo.update({"a":"b"},{$set:{a:"c"}})`
  - changes only one document

# Updating Documents

- Updating multiple documents
  - use the multi option
  - `>db.foo.update({"a":"b"},{$set:{a:"c"}},{multi:true})`
- Updating a document is atomic
  - if two updates happen at the same time, whichever one reaches the server first will be applied, and then the next one will be applied
  - conflicting updates can safely be sent in rapid-fire succession without any documents being corrupted: the last update will “win.”

# Updating Documents

- \$inc to update a key
  - {\$inc: {"score":50}}
  - {\$inc: {"score":500}}
- Array modifiers
  - \$push adds elements to the end of a array
    - create a new array if array does not exist
  - \$ne or \$addToSet
  - \$pop or \$pull to remove elements

# Replacing the document

- To replace the entire content of a document except for the `_id` field, pass an entirely new document as the second argument to `update()`.
  - `>db.foo.update({"a":"c"}, {"b":"d"})`
- If no document matches the update query the `update()` method does nothing
- `{upsert:true}`
  - either update or inserts



# Querying

- `find()`
  - matches every document in the collection and returns them
  - add key/value pair to restrict
  - `> db.users.find({"age" : 27})`
  - `> db.users.find({"username" : "joe"})`
  - `> db.users.find({"username" : "joe", "age" : 27})`

# Specifying which keys to return

- If you do not need all of the key/value pairs, a second argument can be passed specifying the keys
- `> db.users.find({}, {"username" : 1, "email" : 1})`
- `> db.users.find({}, {"username" : 1, "_id" : 0})`

# Query Conditionals

- "\$lt", "\$lte", "\$gt", and "\$gte" are all comparison operators, corresponding to <, <=, >, and >=, respectively.
  - > `db.users.find({"age" : {"$gte" : 18, "$lte" : 30}})`
    - multiple conditionals can be put with a single key
  - > `start = new Date("01/01/2007")`
  - > `db.users.find({"registered" : {"$lt" : start}})`
- \$ne not equal

# OR queries

- \$in can be used to query for a variety of values for a single key.
- allows you to specify criteria of different types as well as values
  - > db.raffle.find({"ticket\_no" : {"\$in" : [725, 542, 390]}})
  - > db.users.find({"user\_id" : {"\$in" : [12345, "joe"]}})
- opposite of \$in is \$nin

# OR queries

- \$or allows to query across different key/value pairs
  - > db.raffle.find({"\$or" : [{"ticket\_no" : 725}, {"winner" : true}]})
  - > db.raffle.find({"\$or" : [{"ticket\_no" : {"\$in" : [725, 542, 390]}}, {"winner" : **true**}]})
- \$not can be applied on top of any other criteria
  - > db.users.find({"id\_num" : {"\$mod" : [5, 1]}})

# Type specific queries

- null
  - Matches with itself but also matches “does not exist”

```
> db.c.find()
```

```
{ "_id" : ObjectId("4ba0f0dfd22aa494fd523621"), "y" : null }
```

```
{ "_id" : ObjectId("4ba0f0dfd22aa494fd523622"), "y" : 1 }
```

```
{ "_id" : ObjectId("4ba0f148d22aa494fd523623"), "y" : 2 }
```

- > db.c.find({"y" : **null**})
- > db.c.find({"z" : null})
- > db.c.find({"z" : {"\$in" : [null], "\$exists" : true}})

# Regular Expressions

- Regular expressions are useful for flexible string matching.
  - For example, if we want to find all users with the name Joe or joe, we can use a regular expression to do caseinsensitive matching:
    - `> db.users.find({"name" : /joe/i})`
    - `> db.users.find({"name" : /joey?/i})`
- MongoDB uses the Perl Compatible Regular Expression (PCRE) library to match regular expressions
- Any regular expression syntax allowed by PCRE is allowed in MongoDB.

# Querying Arrays

- Querying for elements of an array is designed to behave the way querying for scalars
  - `> db.food.insert({"fruit" : ["apple", "banana", "peach"]})`
  - `> db.food.find({"fruit" : "banana"})`
    - Will match the document
- `$all`: allows to match a list of elements
  - `> db.food.find({fruit : {$all : ["apple", "banana"]}})`
    - order does not matter
  - How about
    - `> db.food.find({"fruit" : ["apple", "banana", "peach"]})`
    - `> db.food.find({fruit : ["apple", "banana"]})`
    - `> db.food.find({fruit : ["peach", "apple", "banana"]})`



# Querying Arrays

- *key.index*: allows query for a specific element of an array:
  - `> db.food.find({"fruit.2" : "peach"})`
  - Arrays are always 0-indexed
- *\$size*: allows query for arrays of a given size
  - `> db.food.find({"fruit" : {"$size" : 3}})`
  - *\$size* can not be combined with another conditional say for getting a range of sizes
- *\$slice* can be used to return a subset of elements for an array key
  - `> db.blog.posts.findOne(criteria, {"comments" : {"$slice" : 10}})`
    - return 10 entries.
- range queries
  - `> db.test.find({"x" : {"$gt" : 10, "$lt" : 20}})`
    - `{"x" : 5}`
    - `{"x" : 15}`
    - `{"x" : 25}`
    - `{"x" : [5, 25]}`
  - What will be the output?

# Querying Embedded Documents

- Querying for the whole document: works identically to a normal query

```
{  
  "name" : {  
    "first" : "Joe",  
    "last" : "Schmoe"},  
  "age" : 45  
}
```

- > db.people.find({"name" : {"first" : "Joe", "last" : "Schmoe"}})
- Querying for a specific key or keys of the document
  - > db.people.find({"name.first" : "Joe", "name.last" : "Schmoe"})
  - better than querying the whole document

# Cursors

```
> for(i=0; i<100; i++) {  
... db.collection.insert({x : i});  
... }  
  
> var cursor = db.collection.find();
```

- Cursors generally allow you to control a great deal about the eventual output of a query

```
> while (cursor.hasNext()) {  
... obj = cursor.next();  
... // do stuff  
... }
```

# Limits, Skips, and Sorts

- > `db.c.find().limit(3)`
  - If there are fewer than three documents matching your query in the collection, only the number of matching documents will be returned
  - limit sets an upper limit, not a lower limit.
- > `db.c.find().skip(3)`
  - This will skip the first three matching documents and return the rest of the matches.
  - If there are fewer than three documents in your collection, it will not return any documents.
- > `db.c.find().sort({username : 1, age : -1})`
  - sort takes an object: a set of key/value pairs where the keys are key names and the values are the sort directions.
  - Sort direction can be 1 (ascending) or -1 (descending).
  - If multiple keys are given, the results will be sorted in that order.
- > `db.stock.find({"desc" : "mp3"}).limit(50).sort({"price" : -1})`
- > `db.stock.find({"desc" : "mp3"}).limit(50).skip(50).sort({"price" : -1})`

# Aggregation

- Aggregations are operations that process data records and return computed results
- Aggregation operations in MongoDB use *collections* of documents as an input and return results in the form of one or more documents.
- Single Purpose Aggregation Operations
  - returning a count of matching documents, returning the distinct values for a field, and grouping data based on the values of a field etc.
- MongoDB also provides *map-reduce* operations to perform aggregation

# Single Purpose Aggregation Operations

- **Count:** MongoDB can return a count of the number of documents that match a query.
  - `db.records.count()`
  - `db.records.count( { a: 1 } )`
- **Distinct:** The distinct operation takes a number of documents that match a query and returns all of the unique values for a field in the matching documents.
- **Group :** The *group* operation takes a number of documents that match a query, and then collects groups of documents based on the value of a field or fields.

# Indexing

- To support the efficient execution of queries in MongoDB.
  - Without indexes, MongoDB must perform a collection scan, i.e. scan every document in a collection, to select those documents that match the query statement.
- Indexes are special data structures that store a small portion of the collection's data set in an easy to traverse form
- The index stores the value of a specific field or set of fields, ordered by the value of the field.
- The ordering of the index entries supports efficient equality matches and range-based query operations.
  - In addition, MongoDB can return sorted results by using the ordering in the index.
- MongoDB defines indexes at the *collection* level and supports indexes on any field or sub-field of the documents in a MongoDB collection.

# Index Creation

- `db.collection.createIndex(keys, options)`
- **Keys** : A document that contains the field and value pairs where the field is the index key and the value describes the type of index for that field.
  - For an ascending index on a field, specify a value of 1; for descending index, specify a value of -1.
  - The order of an index is important for supporting `sort()` operations using the index.
- **Options**: Optional. A document that contains a set of options that controls the creation of the index.
  - Unique, sparse, expireAfterSeconds, name, v...



# Index Examples

- `db.collection.createIndex( { orderDate: 1 } )`
  - an ascending index on the field `orderDate`.
- `db.collection.createIndex( { orderDate: 1, zipcode: -1 } )`
- `db.products.createIndex( { item: 1, quantity: -1 } , { name: "inventory" } )`
  - To view the name of an index, use the `getIndexes()` method.

# createIndex() Behaviours

- To add or change index options you must drop the index using the `dropIndex()` method and issue another `createIndex()` operation with the new options.
  - If you create an index with one set of options, and then issue the `createIndex()` method with the same index fields and different options without first dropping the index, `createIndex()` will not rebuild the existing index with the new options.
  - If you call multiple `createIndex()` methods with the same index specification at the same time, only the first operation will succeed, all other operations will have no effect.
- Non-background indexing operations will block all other operations on a database.
  - `db.people.createIndex( { zipcode: 1}, {background: true} )`
- MongoDB will not create an index on a collection if the index entry for an existing document exceeds the Maximum Index Key Length.

# Index Types

- Default `_id`
  - All MongoDB collections have an index on the `_id` field that exists by default.
- Single Field
  - MongoDB supports the creation of user-defined ascending/descending indexes on a single field of a document.
- Compound Index
  - MongoDB also supports user-defined indexes on multiple fields

# Index Types

- Multikey Index
  - MongoDB uses multikey indexes to index the content stored in arrays.
  - If you index a field that holds an array value, MongoDB creates separate index entries for every element of the array.
  - Multikey indexes allow queries to select documents that contain arrays by matching on element or elements of the arrays.

# Other Index types

- Geospatial Index
- Text Index
- Hashed Index

# Index Properties

- TTL Indexes
  - The TTL index is used for TTL collections, which expire data after a period of time.
- Unique Indexes
  - A unique index causes MongoDB to reject all documents that contain a duplicate value for the indexed field.
- Sparse Indexes
  - A sparse index does not index documents that do not have the indexed field.

# TTL Index

- Special single-field indexes that MongoDB can use to automatically remove documents from a collection after a certain amount of time.
  - Data expiration is useful for certain types of information like machine generated event data, logs, and session information
  - `db.eventlog.createIndex( { "lastModifiedDate": 1 }, { expireAfterSeconds: 3600 } )`
  - Expiration of Data
    - TTL indexes expire documents after the specified number of seconds has passed since the indexed field value; i.e. the expiration threshold is the indexed field value plus the specified number of seconds.
  - Restrictions
    - TTL indexes are a single-field indexes. Compound indexes do not support TTL and ignores the `expireAfterSeconds` option.
    - The `_id` field does not support TTL indexes

# Unique Indexes

- A unique index causes MongoDB to reject all documents that contain a duplicate value for the indexed field.
  - *db.members.createIndex( { "user\_id": 1 }, { unique: true } )*
- The unique constraint applies to separate documents in the collection.
- If a document does not have a value for the indexed field in a unique index, the index will store a null value for this document.



# Sparse Indexes

- Sparse indexes only contain entries for documents that have the indexed field, even if the index field contains a null value.
- The index skips over any document that is missing the indexed field.
- The index is “sparse” because it does not include all documents of a collection.
  - By contrast, non-sparse indexes contain all documents in a collection, storing null values for those documents that do not contain the indexed field.
  - `db.addresses.createIndex( { "xmpp_id": 1 }, { sparse: true } )`
- Sparse *compound indexes* that only contain ascending/descending index keys will index a document as long as the document contains at least one of the keys.

# Sparse Indexes

```
{ "_id" : ObjectId("523b6e32fb408eea0eec2647"), "userid" : "newbie" }
```

```
{ "_id" : ObjectId("523b6e61fb408eea0eec2648"), "userid" : "abby", "score" : 82 }
```

```
{ "_id" : ObjectId("523b6e6ffb408eea0eec2649"), "userid" : "nina", "score" : 90 }
```

```
db.scores.createIndex( { score: 1 } , { sparse: true } )
```

```
db.scores.find().sort( { score: -1 } )
```

# Index Management

- To remove an index, use the *db.collection.dropIndex()* method
  - `db.accounts.dropIndex( { "tax-id": 1 } )`
    - The operation returns a document with the status of the operation:
    - `{ "nIndexesWas" : 3, "ok" : 1 }`
- To remove *all* indexes
  - `db.collection.dropIndexes()`
- To modify an Index
  - First drop the old index
  - Then create the new index
- To Rebuild Indexes
  - `db.accounts.reIndex()`
- To return a list of Indexes
  - `db.collection.getIndexes()`

# Index Management

- Return Query Plan with `explain()`
    - Use the `db.collection.explain()` or the `cursor.explain()` method in `executionStats` mode to return statistics about the query process
    - Gives details including the index used, the number of documents scanned, and the time the query takes to process in milliseconds.
  - Control Index Use with `hint()`
    - To force MongoDB to use a particular index for a `db.collection.find()` operation, specify the index with the `hint()` method.
- 
- `db.people.find( { name: "John Doe", zipcode: { $gt: "63000" } }).hint( { zipcode: 1 } ).explain("executionStats")`

# Index Strategies

- Create a Single-Key Index if All Queries Use the Same, Single Key
- Create Compound Indexes to Support Several Different Queries
- If an ascending or a descending index is on a single field, the sort operation on the field can be in either direction.
  - If the query planner cannot obtain the sort order from an index, it will sort the results in memory.
- Ensure Indexes Fit in RAM
  - `db.collection.totalIndexSize()` - returns in bytes

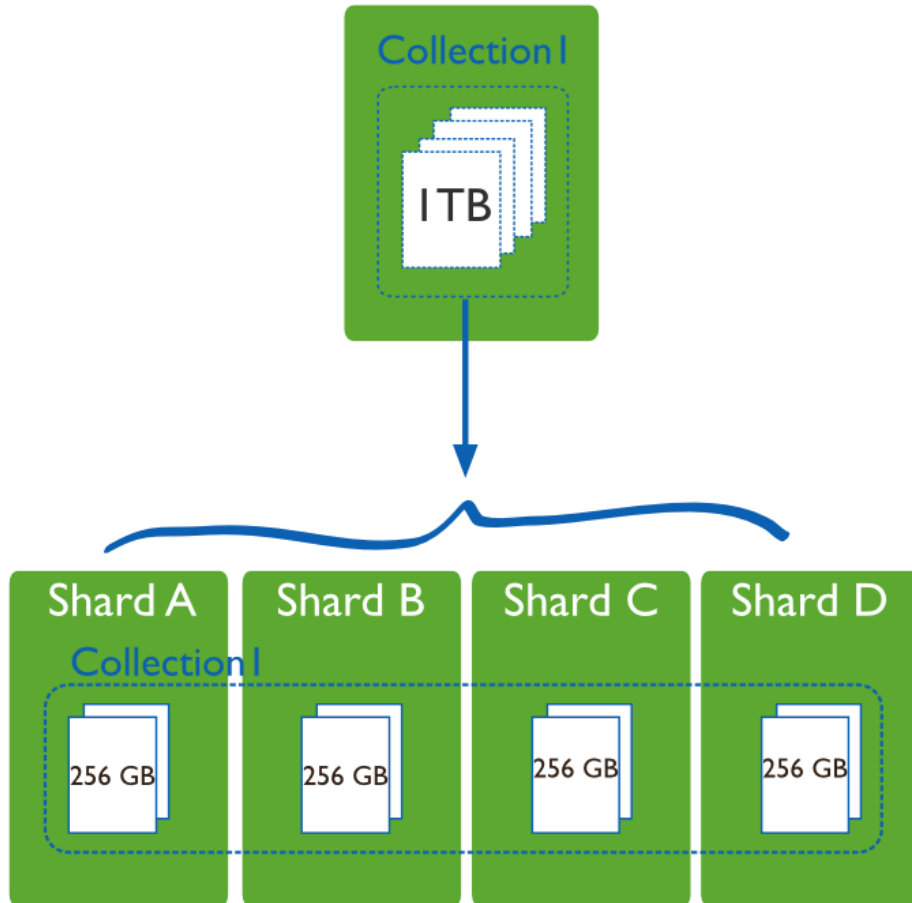
# Indexing considerations

- Each index requires at least 8KB of data space.
- Adding an index has some negative performance impact for write operations. For collections with high write-to-read ratio, indexes are expensive since each insert must also update any indexes.
- Collections with high read-to-write ratio often benefit from additional indexes. Indexes do not affect un-indexed read operations.
- When active, each index consumes disk space and memory. This usage can be significant and should be tracked for capacity planning, especially for concerns over working set size.

# Sharding

- Sharding is a method for storing data across multiple machines.
- Purpose of Sharding
  - Database systems with large data sets and high throughput applications can challenge the capacity of a single server.
  - High query rates can exhaust the CPU capacity of the server.
  - Larger data sets exceed the storage capacity of a single machine.
  - Working set sizes larger than the system's RAM stress the I/O capacity of disk drives.
- Vertical Scaling: adds more CPU and storage resources to increase capacity.
- Horizontal Scaling: divides the data set and distributes the data over multiple servers, or shards

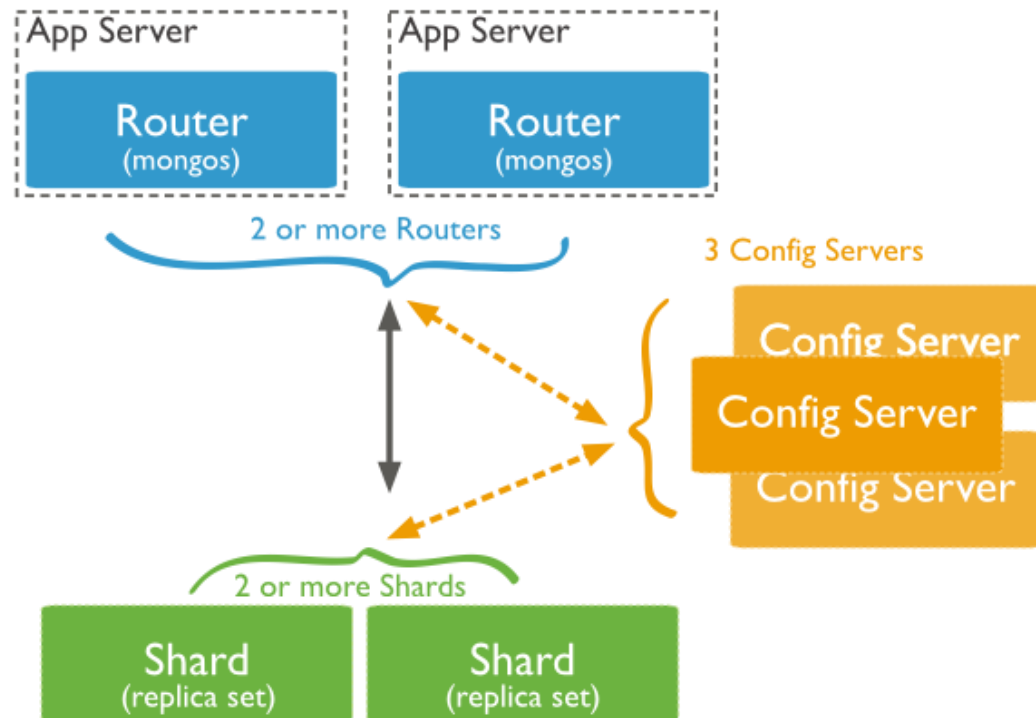
# Sharding



- Sharding reduces the number of operations each shard handles. Each shard processes fewer operations as the cluster grows.
  - As a result, a cluster can increase capacity and throughput horizontally.
- Sharding reduces the amount of data that each server needs to store. Each shard stores less data as the cluster grows.



# Sharding in MongoDB



- Shards store the data. To provide high availability and data consistency, in a production sharded cluster, each shard is a replica set. For more information on replica sets, see Replica Sets.
- Query Routers, or mongos instances, interface with client applications and direct operations to the appropriate shard or shards.
  - The query router processes and targets operations to shards and then returns results to the clients.
  - A sharded cluster can contain more than one query router to divide the client request load.
  - A client sends requests to one query router. Most sharded clusters have many query routers.
- Config servers store the cluster's metadata.
  - The data contains a mapping of the cluster's data set to the shards.
  - The query router uses this metadata to target operations to specific shards. Production sharded clusters have exactly 3 config servers.

# Data Partitioning

- MongoDB distributes data, or shards, at the collection level. Sharding partitions a collection's data by the shard key
- A shard key is either an indexed field or an indexed compound field that exists in every document in the collection.
- MongoDB divides the shard key values into chunks and distributes the chunks evenly across the shards
- Range Based Sharding: MongoDB divides the data set into ranges determined by the shard key values to provide range based partitioning
  - Range based partitioning supports more efficient range queries
- Hash Based Sharding: For hash based partitioning, MongoDB computes a hash of a field's value, and then uses these hashes to create chunks.
  - With hash based partitioning, two documents with "close" shard key values are unlikely to be part of the same chunk. This ensures a more random distribution of a collection in the cluster.